
When Worst-Case Prediction Still Compresses: Tokenization and Adversarial Robustness in LLM-Based Lossless Compression

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 LLM token tables can act as global string compressors: variable-length byte
2 sequences become fixed-width token IDs that support fast decoding and execution
3 on encoded data. Predictive LLM compression adds a second layer, entropy-coding
4 those IDs with next-token probabilities. We show that robustness depends on
5 the product of both layers—tokens per source byte and coded bits per token—
6 and cannot be inferred from either token-table compression ratio or perplexity
7 alone. We construct canonical adversarial sequences and separately attribute
8 tokenization gain, model surprise, arithmetic-coder loss, dictionary transfer, and
9 framing. For Qwen2.5-0.5B, full-vocabulary adversarial sequences cost 5.73 times
10 more bits/token than matched text8, but tokenization absorbs enough of that cost
11 to retain a 51.3% ideal byte rate. When restricted to one-byte tokens, however,
12 the same predictive layer expands data to a 143.4% arithmetic payload. Post-
13 hoc vocabulary masking reduces adversarial model cost by 35.2%. These results
14 identify tokenizer fertility as a robustness boundary and motivate co-designing
15 token tables, predictive coders, and guarded fallback paths.

16 1 What is doing the compression?

17 “Language Modeling Is Compression” formalizes how next-symbol probabilities can drive a lossless
18 arithmetic coder and describes tokenization itself as lossless pre-compression (Delétang et al., 2024).
19 This observation implies that so-called LLM compression combines two distinct mechanisms: a
20 tokenizer replaces variable-length source strings with token IDs, and a language model predicts those
21 IDs so that an entropy coder can shorten likely ones. LLMZip is another instance of this composed
22 design (Valmeekam et al., 2023).

23 Schmidt et al. (Schmidt et al., 2026) isolate the first mechanism. They use an LLM vocabulary as
24 a global symbol table in which common byte strings map to fixed-width IDs and one-byte entries
25 guarantee coverage. Unlike a block-local code, this stable representation supports equality checks,
26 hashing, joins, and aggregations before decoding. These properties make token tables particularly
27 attractive for databases, where strings can remain encoded during common relational operations and
28 be decoded only when needed. Thus tokenization can be a useful compression primitive without
29 language-model inference.

30 What remains unclear is how much of an LLM codec’s behavior comes from this token representation
31 and how much comes from prediction. Delétang et al. compare ASCII and BPE on natural enwik9,
32 but the two effects have not been isolated under inputs chosen against realized end-to-end size.
33 This matters because prediction loss is measured per token while compression is measured per
34 source byte: a poor probability model may still sit inside a pipeline that compresses when its tokens

represent many bytes, whereas low-fertility tokenization can expose expansion. We therefore study the tokenizer–predictor–coder pipeline as a two-layer system.

Contributions. We (i) factor end-to-end rate into tokenizer fertility and predictive coding cost; (ii) construct lossless attacks against token-, byte-, and realized-code objectives; and (iii) use matched natural and adversarial inputs to show when the token layer absorbs predictive surprise and when it exposes expansion.

Scope. We hold one pretrained tokenizer/model fixed and vary canonical inputs adversarially. Our evidence is a Qwen2.5-0.5B case study and does not contradict the token-table paper’s corpus-level results.

2 A two-layer compression model

Let s be UTF-8 text, $T(s) = x_{1:N}$ its tokens, $\text{decode}(x)$ its decoded byte string, $h_t = x_{1:t-1}$, and $p_\theta(v \mid h_t)$ the model distribution. If the codec retains dictionary $A \subseteq V$, it renormalizes that distribution to $q_\theta(v \mid h_t, A)$. The realized sequence NLL is

$$L_A(x) = - \sum_{t=2}^N \log_2 q_\theta(x_t \mid h_t, A). \quad (1)$$

This is sequence cross-entropy, not Shannon entropy of the predictive distribution. By the source-coding correspondence, an ideal arithmetic code has length approximately equal to this NLL (Shannon, 1948; Witten et al., 1987; Delétang et al., 2024).

2.1 Factorizing the two layers

Consider a sequence of N tokens that decodes to U source bytes. Its *fertility*, $f(x) = U/N$, is the average number of bytes represented by one token. A fixed-width token table spends $w = \lceil \log_2 |V| \rceil$ bits per token, or Nw bits in total. Relative to the $8U$ bits in the original bytes, its rate is $R_T = Nw/(8U) = w/(8f)$. The predictive coder replaces the fixed cost w with an average cost $\bar{\ell}_A = L_A/(N-1)$ bits per predicted token. Its ideal rate is therefore

$$R_{\text{pred}} = \frac{L_A}{8U} \approx \frac{\bar{\ell}_A}{8f(x)} = \underbrace{\frac{N}{U}}_{\text{tokens/source byte}} \underbrace{\frac{\bar{\ell}_A}{8}}_{\text{coded bytes/token}}. \quad (2)$$

The predictor improves the token table when $\bar{\ell}_A < w$ and makes its token stream larger when $\bar{\ell}_A > w$. Yet the end-to-end output can still compress if fertility f is large enough. This separates the token-table paper question—how many fixed-width IDs represent the bytes?—from ours—how costly are those IDs under the model?

For arithmetic payload C_{AC} , dictionary cost C_A , and header cost C_H , the realized file rate is

$$R_{\text{file}} = \frac{C_{AC} + C_A + C_H}{8U}. \quad (3)$$

The gaps from Eq. 2 attribute frequency quantization/coder, dictionary, and framing overhead. Values above one denote expansion. We report the raw offline rate and treat model weights as shared decoder state; an adjusted two-part code would additionally charge their size (Delétang et al., 2024).

Compression-oriented objectives. Worst-case prediction loss is not worst-case compression: tokens represent unequal amounts of source data. For actual serialized length C_{actual} in bits, the end-to-end adversary is

$$\max_x \frac{C_{\text{actual}}(x)}{8|\text{decode}(x)|_{\text{bytes}}}. \quad (4)$$

Even a terrible probability model can therefore belong to a pipeline that still compresses because variable-length tokenization is favorable. A cheap greedy approximation chooses high-surprisal tokens representing few decoded bytes:

$$x_t = \arg \max_{v \in V} \frac{-\log_2 p_\theta(v \mid x_{<t})}{|\text{decode}(v)|_{\text{bytes}}}. \quad (5)$$

71 Because standalone decoding can be context-dependent, our implementation uses $\Delta_b(v; h) =$
72 $|\text{decode}(h \parallel v)|_{\text{UTF8}} - |\text{decode}(h)|_{\text{UTF8}}$ in the denominator. Thus minimum probability and maxi-
73 mum surprisal per decoded byte can rank candidates differently.

74 3 Threat model and attacks

75 **Adversary.** Use a white-box adversary with access to tokenizer, logits, and coding algorithm. The
76 adversary chooses a length- N token sequence after a fixed seed, subject to a candidate alphabet
77 G . Because arbitrary token IDs can change after decoding and retokenizing, we accept a prefix
78 only when $T(\text{decode}(x_{1:t})) = x_{1:t}$; final decoding must recover the same IDs and bytes exactly.
79 We evaluate the full non-special vocabulary, printable ASCII, and canonical one-byte ASCII. The
80 last isolates predictive failure from token-length variation; broader alphabets expose the tokenizer
81 interface.

82 Attack variants.

Variant	How the next token or sequence is chosen
Random token	Seeded uniform sampling among valid extensions; no model-based objective.
83 Minimum probability	Greedily choose the valid token with smallest $p_\theta(v \mid x_{<t})$.
Surprisal/byte	Greedily maximize $-\log_2 p_\theta(v \mid x_{<t}) / \Delta_b(v; x_{<t})$ (Eq. 5).
Serialized-size search	Beam search that clones coder state and maximizes realized bits per decoded byte (Eq. 4).

84 Controls are ordinary text8, seeded random valid UTF-8, and unrestricted random bytes for byte-
85 capable codecs. We apply gzip, Zstandard, and Brotli to every exact decoded sequence. Beam search
86 is approximate unless its branch factor covers all candidates; changing beam ancestry disables the
87 KV cache.

88 **Global mask.** We construct each adversarial sequence once under the full distribution, then replay
89 identical IDs and histories under its occurring-token mask. This isolates renormalization and is not a
90 mask-adaptive attack.

91 4 Experimental evaluation

92 **Setup.** We evaluate Qwen2.5-0.5B with its 151,643-token vocabulary. The adversary creates three
93 1,000-token sequences from starting token IDs 785, 32, and 641 by repeatedly choosing the lowest-
94 probability canonical extension. Every completed sequence satisfies $T(\text{decode}(x)) = x$ and decodes
95 to 6,988, 6,812, and 6,943 UTF-8 bytes. We then replay the identical token IDs and contexts using
96 the set of tokens occurring in each completed sequence; the three post-hoc dictionaries contain 592,
97 606, and 587 tokens. Thus masking changes only probability normalization, not generation.

98 The natural-text control is the first 100,000 text8 tokens, processed by the same model with context
99 length 1,000, retained context 100, KV caching, and arithmetic coding. It uses 16 independently
100 seeded batches, so 99,984 tokens are predicted. We report mean adversarial surprisal with sample
101 standard deviation across the three starts. Text8 is one deterministic matched run per dictionary policy
102 and therefore has no error bar.

103 **Token-level failure, byte-level buffering.** Figure 1 shows that full-vocabulary adversarial inputs
104 cost 28.409 ± 0.077 ideal bits per predicted token, $5.734 \times$ text8 at 4.954. Because Qwen has
105 151,643 symbols, a fixed ID needs 18 bits; the predictive layer therefore makes this adversarial token
106 stream 58% larger than fixed-width IDs. However, the sequences decode to an average 6.914 source
107 bytes/token. This strong token-table compression buffers model surprise, leaving an ideal end-to-end
108 rate of 51.346% of raw UTF-8 rather than expansion.

109 Post-hoc occurring-token dictionaries contain 592, 606, and 587 symbols. Masking the same IDs
110 and histories lowers adversarial cost to 18.402 ± 0.833 bits/token, a 35.2% reduction, and the byte-
111 normalized floor to 33.284%. It remains $3.806 \times$ the matched masked-text8 cost of 4.835 bits/token.
112 The identical full-model scores under both policies verify that only normalization changes.




Figure 1: How tokenization and predictive coding jointly determine size for Qwen2.5-0.5B. (a) Fixed-width token-table rate, model-induced floor, and realized coding costs relative to each input’s UTF-8 bytes. The full-vocabulary adversarial arithmetic payload has not been measured. (b) Post-hoc masking rescoring of identical token IDs and histories; masking reduces the adversarial floor by 35.2%.

113 **Typical text and realized overhead.** For text8, model-induced floors are 11.631% of raw UTF-8
 114 with the full vocabulary and 11.352% with occurring tokens. Finite coding and metadata raise realized
 115 totals to 14.405% and 13.434%, or $6.942\times$ and $7.444\times$ compression. The occurring-token bitmap
 116 costs 10,818 bytes (2.03% of raw size), less than its coding gain at 100,000 tokens.

117 **Removing tokenization gain exposes expansion.** The notebook one-byte-ASCII test fixes fertility
 118 at one byte/token. Greedy minimum probability and surprisal-per-byte then coincide at a 135.48%
 119 model floor and 143.36% arithmetic payload; an NLL beam reaches 136.70% and 144.14%. Thus
 120 predictive coding improves on 18-bit fixed Qwen IDs but cannot beat the original 8-bit bytes. The
 121 narrow 32-token beam validates this mechanism and the accounting, not global optimality. Together,
 122 the experiments show both regimes of Eq. 2: multi-byte tokens can absorb extreme surprise, while
 123 low-fertility inputs expose predictive expansion.

124 5 Systems implications and limitations

125 The systems lesson is to preserve token IDs as the global, queryable in-memory representation
 126 and treat predictive arithmetic coding as an optional block-level storage or transport layer. Source-
 127 normalized accounting identifies whether fragility comes from the tokenizer, predictor, quantizer,
 128 dictionary, or container. For untrusted blocks, a guard can store the smaller of predictive and
 129 conventional encodings plus a selector bit, bounding size by the fallback plus one bit; its decoding
 130 granularity, latency, and throughput remain to be measured.

131 Our evidence is limited to one model, synthetic white-box inputs, three starts, finite context, ap-
 132 proximate search, and one coder/container. It does not establish global worst cases or generalize
 133 to multilingual text or code. The attack could enable storage or compute amplification; practical
 134 mitigations include size caps, early-abort expansion detection, bounded inference, and fallback
 135 codecs.

136 6 Conclusion

137 Robust learned compression is an end-to-end systems problem. Canonical byte-aware attacks and
 138 explicit cost attribution show that the tested adversarial sequences are 3.81–5.73 times harder than
 139 matched text8, while guarded fallback offers a principled path from diagnosis to robust operation.

References

- Grégoire Delétang, Anian Ruoss, Paul-Ambroise Duquenne, Elliot Catt, Tim Genewein, Christopher Mattern, Jordi Grau-Moya, Li Kevin Wenliang, Matthew Aitchison, Laurent Orseau, Marcus Hutter, and Joel Veness. Language modeling is compression. In *ICLR*, 2024.
- Tobias Schmidt, Nicolas Schmitt, Thomas Neumann, and Andreas Kipf. Accelerating string-heavy queries with LLM token tables. Manuscript, 2026.
- Claude E. Shannon. A mathematical theory of communication. *Bell System Technical Journal*, 27:379–423, 623–656, 1948.
- Chandra S. K. Valmeekam, Krishna Narayanan, Dileep Kalathil, Jean-François Chamberland, and Srinivas Shakkottai. LLMZip: Lossless text compression using large language models. arXiv:2306.04050, 2023.
- Ian H. Witten, Radford M. Neal, and John G. Cleary. Arithmetic coding for data compression. *Communications of the ACM*, 30(6):520–540, 1987.

NeurIPS Paper Checklist

The checklist is designed to encourage best practices for responsible machine learning research, addressing issues of reproducibility, transparency, research ethics, and societal impact. Do not remove the checklist: **The papers not including the checklist will be desk rejected.** The checklist should follow the references and follow the (optional) supplemental material. The checklist does NOT count towards the page limit.

Please read the checklist guidelines carefully for information on how to answer these questions. For each question in the checklist:

- You should answer [Yes], [No], or [N/A].
- [N/A] means either that the question is Not Applicable for that particular paper or the relevant information is Not Available.
- Please provide a short (1–2 sentence) justification right after your answer (even for [N/A]).

The checklist answers are an integral part of your paper submission. They are visible to the reviewers, area chairs, senior area chairs, and ethics reviewers. You will also be asked to include it (after eventual revisions) with the final version of your paper, and its final version will be published with the paper.

The reviewers of your paper will be asked to use the checklist as one of the factors in their evaluation. While [Yes] is generally preferable to [No], it is perfectly acceptable to answer [No] provided a proper justification is given (e.g., error bars are not reported because it would be too computationally expensive” or “we were unable to find the license for the dataset we used”). In general, answering [No] or [N/A] is not grounds for rejection. While the questions are phrased in a binary way, we acknowledge that the true answer is often more nuanced, so please just use your best judgment and write a justification to elaborate. All supporting evidence can appear either in the main paper or the supplemental material, provided in appendix. If you answer [Yes] to a question, in the justification please point to the section(s) where related material for the question can be found.

IMPORTANT, please:

- **Delete this instruction block, but keep the section heading “NeurIPS Paper Checklist”.**
- **Keep the checklist subsection headings, questions/answers and guidelines below.**
- **Do not modify the questions and only use the provided macros for your answers.**

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that the paper does not include experiments.

- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).

- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer **[N/A]** means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer **[N/A]** means that the paper does not include experiments.
- The authors should answer **[Yes]** if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer **[N/A]** means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.

- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines?>

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer **[No]**, they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or **[No]**, they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that the paper poses no such risks.

- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. Declaration of LLM usage

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.